

Experiment Results Summary
(Alibaba 2026 GenAI Trace, Load=8 jobs/hr, 128 GPUs)

Scheduler	Avg JCT (h)	vs FIFO	Median (h)	P95 (h)	P99 (h)	Resp (s)	SLO@2x (%)
FIFO	29.2	baseline	19.5	75.5	134.8	52,573	23.8
LAS	14.9	-49.0%	4.8	60.4	126.3	456	100.0
SRTF	14.9	-49.2%	4.8	60.4	121.0	510	100.0
SloScoring	17.6	-39.8%	7.6	65.2	124.8	8,529	64.4